

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler,LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, roc_curve, auc
import seaborn as sns

"""For DT plotting"""
from io import StringIO
from IPython.display import Image
from sklearn.tree import export_graphviz
import pydotplus
```

```
In [ ]:
```

```
In [2]: retail = pd.read_csv("Retail_Customer_Insights.csv")
retail.head()
```

Out[2]:

	Customer_ID	Age	Annual_Income	Gender	Purchase_History	Product_Category	Customer_Satisfaction	Loyalty_Points	Marital
0	CID770487	45	72633.53	Non-binary	0	Electronics	9.0	541.11	I
1	CID216739	38	61816.55	Non-binary	0	Books	6.0	497.41	
2	CID126225	47	57338.15	Non-binary	0	Grocery	3.0	634.90	
3	CID877572	58	83800.37	Female	0	Furniture	4.0	505.82	I
4	CID388389	37	64875.12	Male	0	Furniture	6.0	610.39	

```
In [3]: retail.isnull().sum()
```

Out[3]:

Customer_ID	0
Age	0
Annual_Income	5000
Gender	0
Purchase_History	0
Product_Category	0
Customer_Satisfaction	3000
Loyalty_Points	2000
Marital_Status	0
Number_of_Children	0
Employment_Status	0
Credit_Score	0
Owns_House	0
Monthly_Expenditure	5000
Internet_Usage_Hours_per_Week	0
dtype:	int64

```
In [4]: retail.isna().sum()
```

Out[4]:

Customer_ID	0
Age	0
Annual_Income	5000
Gender	0
Purchase_History	0
Product_Category	0
Customer_Satisfaction	3000
Loyalty_Points	2000
Marital_Status	0
Number_of_Children	0
Employment_Status	0
Credit_Score	0
Owns_House	0
Monthly_Expenditure	5000
Internet_Usage_Hours_per_Week	0
dtype:	int64

```
In [5]: retail.fillna({'Annual_Income':retail['Annual_Income'].median(),
                    'Customer_Satisfaction':retail['Customer_Satisfaction'].median(),
                    'Loyalty_Points':retail['Loyalty_Points'].median(),
                    'Monthly_Expenditure':retail['Monthly_Expenditure'].median()},inplace=True)
```

```
In [6]: retail.isna().sum()
```

```
Out[6]: Customer_ID      0
        Age              0
        Annual_Income    0
        Gender           0
        Purchase_History 0
        Product_Category 0
        Customer_Satisfaction 0
        Loyalty_Points    0
        Marital_Status    0
        Number_of_Children 0
        Employment_Status 0
        Credit_Score      0
        Owns_House        0
        Monthly_Expenditure 0
        Internet_Usage_Hours_per_Week 0
        dtype: int64
```

```
In [7]: retail.isnull().sum()
```

```
Out[7]: Customer_ID      0
        Age              0
        Annual_Income    0
        Gender           0
        Purchase_History 0
        Product_Category 0
        Customer_Satisfaction 0
        Loyalty_Points    0
        Marital_Status    0
        Number_of_Children 0
        Employment_Status 0
        Credit_Score      0
        Owns_House        0
        Monthly_Expenditure 0
        Internet_Usage_Hours_per_Week 0
        dtype: int64
```

```
In [8]: # Removing outliers from Age column
dataset=np.array(retail['Age']).tolist()
dataset.sort()
median_ = np.median(dataset)
q1=np.percentile(dataset,25)
q3=np.percentile(dataset,75)
iqr=q3-q1
ll=q1-1.5*iqr
ul=q3+1.5*iqr
data=retail.loc[(retail['Age']>ll) & (retail['Age']<ul)]
```

```
In [9]: # Removing outliers from Annual_Income column
dataset=np.array(data['Annual_Income']).tolist()
dataset.sort()
median_ = np.median(dataset)
q1=np.percentile(dataset,25)
q3=np.percentile(dataset,75)
iqr=q3-q1
ll=q1-1.5*iqr
ul=q3+1.5*iqr
data=data.loc[(data['Annual_Income']>ll) & (data['Annual_Income']<ul)]
```

```
In [10]: # Removing outliers from Loyalty_Points column
dataset=np.array(data['Loyalty_Points']).tolist()
dataset.sort()
median_ = np.median(dataset)
q1=np.percentile(dataset,25)
q3=np.percentile(dataset,75)
iqr=q3-q1
ll=q1-1.5*iqr
ul=q3+1.5*iqr
data=data.loc[(data['Loyalty_Points']>ll) & (data['Loyalty_Points']<ul)]
```

```
In [11]: # Removing outliers from Number_of_Children column
dataset=np.array(data['Number_of_Children']).tolist()
dataset.sort()
median_ = np.median(dataset)
q1=np.percentile(dataset,25)
q3=np.percentile(dataset,75)
iqr=q3-q1
ll=q1-1.5*iqr
ul=q3+1.5*iqr
data=data.loc[(data['Number_of_Children']>ll) & (data['Number_of_Children']<ul)]
```

```
In [12]: # Removing outliers from Credit_Score column
dataset=np.array(data['Credit_Score']).tolist()
```

```
dataset.sort()
median_ = np.median(dataset)
q1=np.percentile(dataset,25)
q3=np.percentile(dataset,75)
iqr=q3-q1
ll=q1-1.5*iqr
ul=q3+1.5*iqr
data=data.loc[(data['Credit_Score']>ll) & (data['Credit_Score']<ul)]
```

```
In [13]: # Removing outliers from Number_of_Children column
dataset=np.array(data['Monthly_Expenditure']).tolist()
dataset.sort()
median_ = np.median(dataset)
q1=np.percentile(dataset,25)
q3=np.percentile(dataset,75)
iqr=q3-q1
ll=q1-1.5*iqr
ul=q3+1.5*iqr
data=data.loc[(data['Monthly_Expenditure']>ll) & (data['Monthly_Expenditure']<ul)]
```

```
In [14]: # Removing outliers from Internet_Usage_Hours_per_Week column
dataset=np.array(data['Internet_Usage_Hours_per_Week']).tolist()
dataset.sort()
median_ = np.median(dataset)
q1=np.percentile(dataset,25)
q3=np.percentile(dataset,75)
iqr=q3-q1
ll=q1-1.5*iqr
ul=q3+1.5*iqr
data=data.loc[(data['Internet_Usage_Hours_per_Week']>ll) & (data['Internet_Usage_Hours_per_Week']<ul)]
```

```
In [15]: # setting customer_id as index
data.set_index('Customer_ID',inplace=True)
data.head()
```

```
Out[15]:
```

	Age	Annual_Income	Gender	Purchase_History	Product_Category	Customer_Satisfaction	Loyalty_Points	Marital_St
Customer_ID								
CID770487	45	72633.53	Non-binary	0	Electronics	9.0	541.11	Divc
CID216739	38	61816.55	Non-binary	0	Books	6.0	497.41	Ma
CID126225	47	57338.15	Non-binary	0	Grocery	3.0	634.90	S
CID877572	58	83800.37	Female	0	Furniture	4.0	505.82	Divc
CID356787	37	57270.25	Prefer not to say	1	Grocery	3.0	458.98	Divc

```
In [16]: data.info()

<class 'pandas.core.frame.DataFrame'>
Index: 81315 entries, CID770487 to CID966793
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    81315 non-null  int64
1   Annual_Income                         81315 non-null  float64
2   Gender                                81315 non-null  object
3   Purchase_History                      81315 non-null  int64
4   Product_Category                      81315 non-null  object
5   Customer_Satisfaction                 81315 non-null  float64
6   Loyalty_Points                       81315 non-null  float64
7   Marital_Status                       81315 non-null  object
8   Number_of_Children                   81315 non-null  int64
9   Employment_Status                    81315 non-null  object
10  Credit_Score                          81315 non-null  int64
11  Owns_House                            81315 non-null  bool
12  Monthly_Expenditure                  81315 non-null  float64
13  Internet_Usage_Hours_per_Week        81315 non-null  int64
dtypes: bool(1), float64(4), int64(5), object(4)
memory usage: 8.8+ MB
```

```
In [17]: cat_cols=['Gender','Product_Category','Marital_Status','Employment_Status','Owns_House']
for i in cat_cols:
    le=LabelEncoder()
    data[i]=le.fit_transform(data[i])
```

```
In [18]: data.head()
```

```
Out[18]:
```

	Age	Annual_Income	Gender	Purchase_History	Product_Category	Customer_Satisfaction	Loyalty_Points	Marital_St
Customer_ID								
CID770487	45	72633.53	2	0	2	9.0	541.11	
CID216739	38	61816.55	2	0	0	6.0	497.41	
CID126225	47	57338.15	2	0	4	3.0	634.90	
CID877572	58	83800.37	0	0	3	4.0	505.82	
CID356787	37	57270.25	3	1	4	3.0	458.98	

```
In [19]: data.shape
```

```
Out[19]: (81315, 14)
```

```
In [20]: x=data.drop('Purchase_History',axis=1)
y=data['Purchase_History']
```

```
In [21]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=74)
```

```
In [22]: #Feature Scaling
scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)
```

```
In [23]: model_dt_2 = DecisionTreeClassifier(max_depth=2)
model_dt_2.fit(x_train_scaled,y_train)
model_dt_2_tr_score = model_dt_2.score(x_train_scaled,y_train)
model_dt_2_te_score = model_dt_2.score(x_test_scaled,y_test)

print(f'Training Score: {model_dt_2_tr_score}')
print(f'Test Score: {model_dt_2_te_score}')
```

Training Score: 0.5984904384184959  
Test Score: 0.6006271905552482

```
In [24]: model_dt_4 = DecisionTreeClassifier(max_depth=4)
model_dt_4.fit(x_train_scaled,y_train)
model_dt_4_tr_score = model_dt_4.score(x_train_scaled,y_train)
model_dt_4_te_score = model_dt_4.score(x_test_scaled,y_test)

print(f'Training Score: {model_dt_4_tr_score}')
print(f'Test Score: {model_dt_4_te_score}')
```

Training Score: 0.5987210231814548  
Test Score: 0.6003812334747587

```
In [25]: model_dt_8 = DecisionTreeClassifier(max_depth=8)
model_dt_8.fit(x_train_scaled,y_train)
model_dt_8_tr_score = model_dt_8.score(x_train_scaled,y_train)
model_dt_8_te_score = model_dt_8.score(x_test_scaled,y_test)

print(f'Training Score: {model_dt_8_tr_score}')
print(f'Test Score: {model_dt_8_te_score}')
```

Training Score: 0.6021490499907766  
Test Score: 0.5976142163192523

```
In [26]: model_dt_16 = DecisionTreeClassifier(max_depth=16)
model_dt_16.fit(x_train_scaled,y_train)
model_dt_16_tr_score = model_dt_16.score(x_train_scaled,y_train)
model_dt_16_te_score = model_dt_16.score(x_test_scaled,y_test)

print(f'Training Score: {model_dt_16_tr_score}')
print(f'Test Score: {model_dt_16_te_score}')
```

Training Score: 0.6284049683330258  
Test Score: 0.5891286970423661

```
In [27]: model_dt_2 = DecisionTreeClassifier(max_depth=2,criterion='entropy')
model_dt_2.fit(x_train_scaled,y_train)
model_dt_2_tr_score = model_dt_2.score(x_train_scaled,y_train)
model_dt_2_te_score = model_dt_2.score(x_test_scaled,y_test)

print(f'Training Score: {model_dt_2_tr_score}')
print(f'Test Score: {model_dt_2_te_score}')
```

Training Score: 0.5984904384184959  
Test Score: 0.6006271905552482

```
In [29]: model_dt_4 = DecisionTreeClassifier(max_depth=4,criterion='entropy')
model_dt_4.fit(x_train_scaled,y_train)
model_dt_4_tr_score = model_dt_4.score(x_train_scaled,y_train)
model_dt_4_te_score = model_dt_4.score(x_test_scaled,y_test)

print(f'Training Score: {model_dt_4_tr_score}')
print(f'Test Score: {model_dt_4_te_score}')
```

Training Score: 0.5986595339113324  
Test Score: 0.6002582549345139

```
In [31]: model_dt_8 = DecisionTreeClassifier(max_depth=8,criterion='entropy')
model_dt_8.fit(x_train_scaled,y_train)
model_dt_8_tr_score = model_dt_8.score(x_train_scaled,y_train)
model_dt_8_te_score = model_dt_8.score(x_test_scaled,y_test)

print(f'Training Score: {model_dt_8_tr_score}')
print(f'Test Score: {model_dt_8_te_score}')
```

Training Score: 0.6007194244604317  
Test Score: 0.5985980446412101

```
In [33]: tr_prediction = model_dt_2.predict(x_train_scaled)

print(model_dt_2.score(x_train_scaled,y_train))
```

0.5984904384184959

```
In [35]: te_prediction = model_dt_2.predict(x_test_scaled)

print(model_dt_2.score(x_test_scaled,y_test))
```

0.6006271905552482


```
In [37]: cm_tr = confusion_matrix(y_train,tr_prediction)
cm_tr
```

```
Out[37]: array([[38906,    8],
               [26111,   27]], dtype=int64)
```


```
In [39]: cm_te = confusion_matrix(y_test,te_prediction)
cm_te
```

```
Out[39]: array([[9767,   10],
               [6485,    1]], dtype=int64)
```

```
In [41]: sns.set({'figure.figsize':(3,2)})
sns.heatmap(cm_tr,fmt='d',annot=True,cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix(Test)')
plt.show()
```



```
In [43]: sns.set({'figure.figsize':(3,2)})
sns.heatmap(cm_te,fmt='d',annot=True,cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix(Test)')
plt.show()
```



```
In [45]: from sklearn.ensemble import AdaBoostClassifier
base_classifier = DecisionTreeClassifier(max_depth = 2)
adaboost_classifier = AdaBoostClassifier(estimator = base_classifier,n_estimators = 40, random_state = 37)
adaboost_classifier.fit(x_train_scaled, y_train)
y_train_pred = adaboost_classifier.predict(x_train_scaled)
print(adaboost_classifier.score(x_train_scaled, y_train))
print(adaboost_classifier.score(x_test_scaled, y_test))
```

0.6007194244604317  
0.5974297485088852

```
In [47]: from sklearn.ensemble import AdaBoostClassifier
base_classifier = DecisionTreeClassifier(max_depth = 2)
adaboost_classifier = AdaBoostClassifier(estimator = base_classifier,n_estimators = 30, random_state = 37)
adaboost_classifier.fit(x_train_scaled, y_train)
y_train_pred = adaboost_classifier.predict(x_train_scaled)
print(adaboost_classifier.score(x_train_scaled, y_train))
print(adaboost_classifier.score(x_test_scaled, y_test))
```

0.6004273504273504  
0.5985980446412101

```
In [49]: from sklearn.ensemble import AdaBoostClassifier
base_classifier = DecisionTreeClassifier(max_depth = 1)
adaboost_classifier = AdaBoostClassifier(estimator = base_classifier,n_estimators = 30, random_state = 37)
adaboost_classifier.fit(x_train_scaled, y_train)
y_train_pred = adaboost_classifier.predict(x_train_scaled)
print(adaboost_classifier.score(x_train_scaled, y_train))
print(adaboost_classifier.score(x_test_scaled, y_test))
```

0.5985826723236796  
0.6005042120150034

```
In [51]: from sklearn.ensemble import AdaBoostClassifier
base_classifier = DecisionTreeClassifier(max_depth = 1)
adaboost_classifier = AdaBoostClassifier(estimator = base_classifier,n_estimators = 50, random_state = 37)
adaboost_classifier.fit(x_train_scaled, y_train)
y_train_pred = adaboost_classifier.predict(x_train_scaled)
print(adaboost_classifier.score(x_train_scaled, y_train))
print(adaboost_classifier.score(x_test_scaled, y_test))
```

0.5986595339113324  
0.600442722744881

In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js